



# Advanced Database Systems

Spring 2024

## Lecture #25: Locking

R&G: Chapters 16 & 17

1

## QUERY SCHEDULER

### How to guarantee only serializable schedules in DBMS?

Problem: user does not need to specify the full transaction at once

Goal: build a query scheduler that always emits serializable schedules

#### Pessimistic (locking)

Use locks to protect database objects

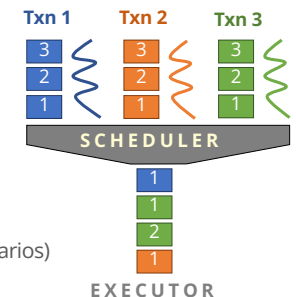
Standard approach if conflicts are frequent

#### Optimistic (versioning)

Record changes for each txn individually

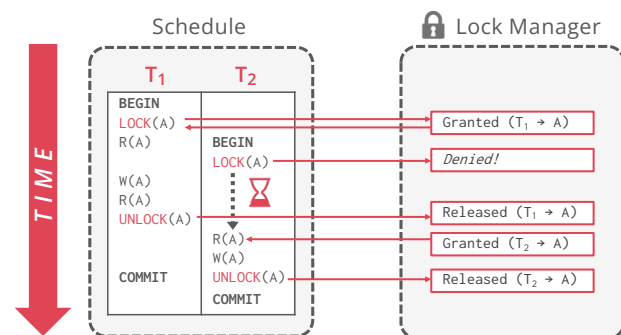
Validate and possibly rollback on commit

Used if conflicts are rare (e.g., write-once-read-many scenarios)



2

## EXECUTING WITH LOCKS



3

## EXECUTING WITH LOCKS

Basic lock types:

**S-LOCK:** Shared locks for reads

**X-LOCK:** Exclusive locks for writes

Steps:

Transactions request locks (or upgrades) before accessing objects

Lock manager grants or blocks requests

Transactions release locks

Lock manager updates its internal lock-table

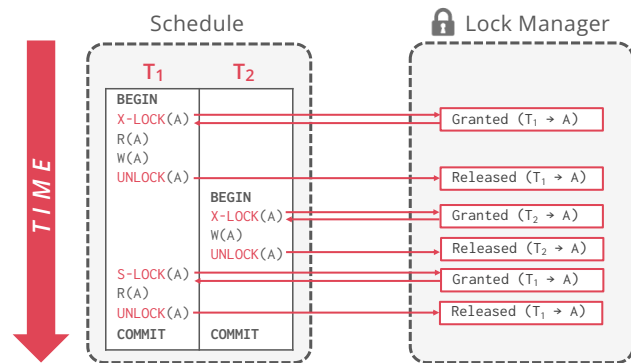
### Compatibility Matrix

|           | Shared | Exclusive |
|-----------|--------|-----------|
| Shared    | ✓      | ✗         |
| Exclusive | ✗      | ✗         |

4

## EXECUTING WITH LOCKS

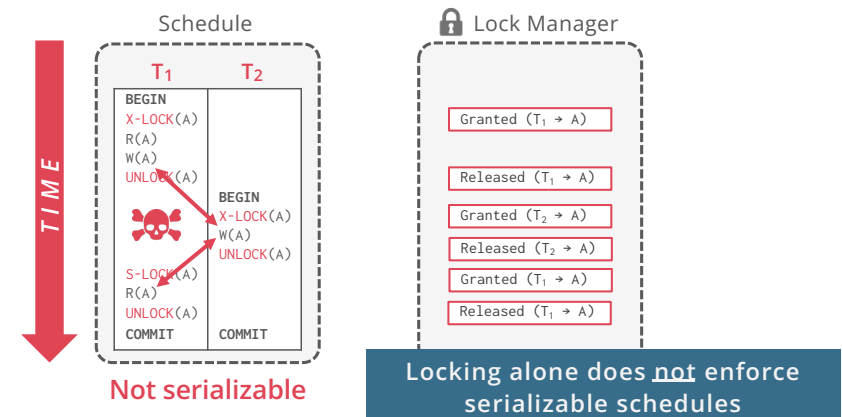
5



5

## EXECUTING WITH LOCKS

6



6

## TWO-PHASE LOCKING

7

### Locks + concurrency control protocol

Determines if a txn is allowed to access an object in the database on the fly  
Does not need to know all of the queries that a txn will execute ahead of time

### Phase 1: Growing

Each txn requests the locks that it needs from the lock manager  
The lock manager grants/denies lock requests

### Phase 2: Shrinking


The txn is allowed to only release locks that it previously acquired  
It cannot acquire new locks

7

## TWO-PHASE LOCKING

8


The transaction is not allowed to acquire/upgrade locks after the growing phase finishes



8

## TWO-PHASE LOCKING

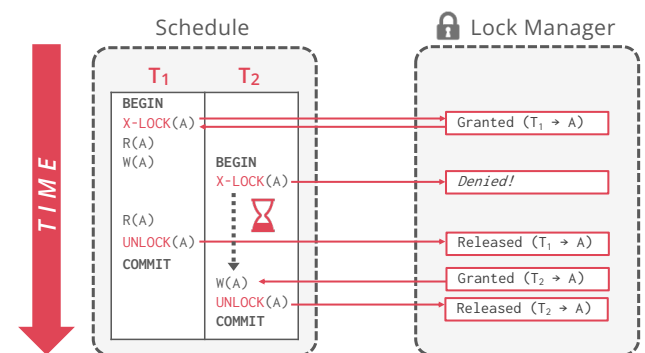
The transaction is not allowed to acquire/upgrade locks after the growing phase finishes



9

9

## EXECUTING WITH LOCKS

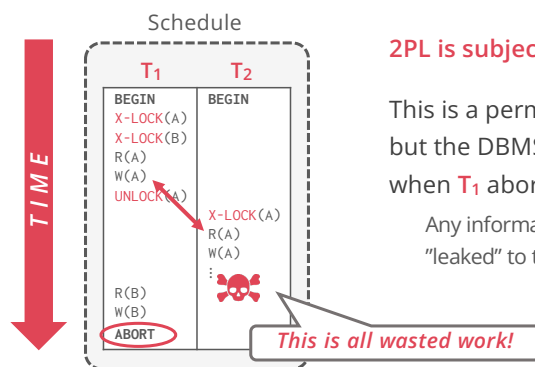


**2PL is sufficient to guarantee conflict-serializability**  
(generates schedules whose precedence graph is acyclic)

10

10

## 2PL – CASCADING ABORTS



**2PL is subject to cascading aborts**

This is a permissible schedule in 2PL  
but the DBMS has to also abort **T<sub>2</sub>**  
when **T<sub>1</sub>** aborts

Any information about **T<sub>1</sub>** cannot be  
"leaked" to the outside world

11

11

## 2PL OBSERVATIONS

There are schedules that are serializable but not be allowed by 2PL

Locking limits concurrency

May require **cascading aborts**

Solution: Strict 2PL

May still have **"dirty reads"**

Solution: Strict 2PL

May lead to **deadlocks**

Solution: Detection or Prevention


12

12

## STRICT TWO-PHASE LOCKING

13

The txn is not allowed to acquire/upgrade locks after the growing phase finishes  
Allows only conflict-serializable schedules, but it is often stronger than needed for some applications



13


## STRICT TWO-PHASE LOCKING

14

Advantages:

Does not incur cascading aborts


Aborted txns can be undone by just restoring original values of modified tuples



14

## NON-2PL EXAMPLE

15



T<sub>1</sub> – move £100 from account A to account B

T<sub>2</sub> – compute the total amount in all accounts and return it to the application

Initial Database State

A = 1000, B = 1000


T<sub>2</sub> Output

A + B = 1900

15

## 2PL EXAMPLE

16



T<sub>1</sub> – move £100 from account A to account B

T<sub>2</sub> – compute the total amount in all accounts and return it to the application

Initial Database State

A = 1000, B = 1000

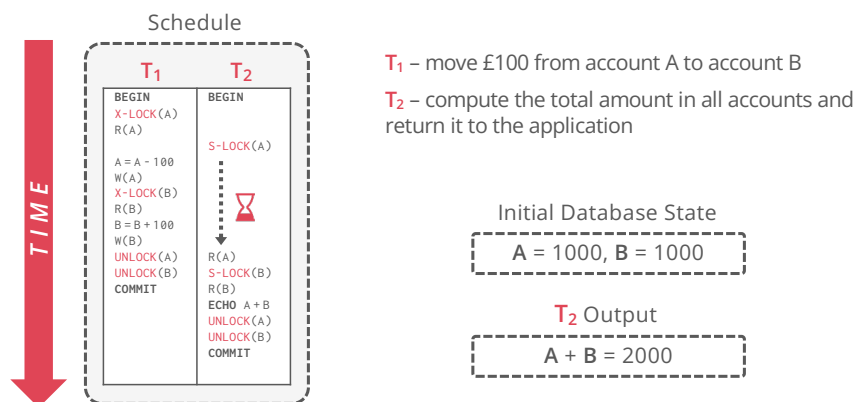
T<sub>2</sub> Output

A + B = 2000

16

## STRICT 2PL EXAMPLE

17

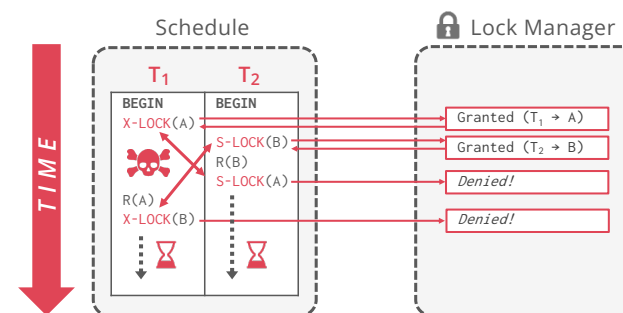


17

## SCHEDULING: DEADLOCKS

18

Two-phase locking has the risk of **deadlock situations**



18

## 2PL DEADLOCKS

19


**Deadlock** = a cycle of txns waiting for locks to be released by each other

Two ways of dealing with deadlocks:

Deadlock Detection

Deadlock Prevention

Conservative (or “preclaiming”) 2PL also prevents deadlocks. Why?



19

Gain all the locks that are necessary before executing the transaction.  
Generally we don't have this info beforehand hence this is not used much

## DEADLOCK DETECTION

20

The DBMS creates a **waits-for** graph to keep track of what locks each transaction is waiting to acquire:

Nodes are transactions

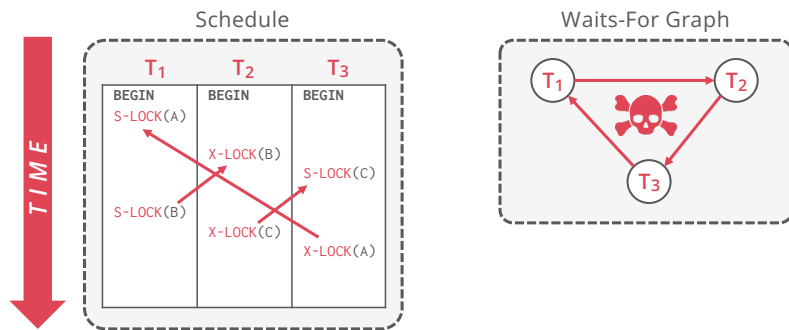
Edge from T<sub>i</sub> to T<sub>j</sub> if T<sub>i</sub> is waiting for T<sub>j</sub> to release a lock

The system periodically checks for cycles in waits-for graph and then make a decision on how to break it

20

## DEADLOCK DETECTION

21



21

## DEADLOCK HANDLING

22

Upon detecting a deadlock, the DBMS selects a “**victim**” transaction to rollback to break the cycle

Selecting a “victim” transaction might depend on:

- age (lowest timestamp)
- progress (least/most executed queries)
- # of items already locked
- # of txns that we have to rollback with it
- # of previous restarts (to prevent starvation)

There is a trade-off between the frequency of checking for deadlocks and how long transactions have to wait before deadlocks are broken

Common frequency of check: 1second

22

## DEADLOCK PREVENTION

23

When a transaction tries to acquire a lock that is held by another transaction, kill one of them to prevent a deadlock

No waits-for graph or detection algorithm

Assign **priorities** based on timestamps

Older  $\Rightarrow$  higher priority (e.g.,  $T_1 > T_2$ )

Two deadlock prevention policies:

**Wait-Die** (“Old Waits for Young”)

**Wound-Wait** (“Young Waits for Old”)

23

## DEADLOCK PREVENTION

24

### Wait-Die (“Old Waits for Young”)

If *requesting* txn has higher priority than *holding* txn

Then *requesting* txn **waits** for *holding* txn

Else *requesting* txn **aborts**


### Wound-Wait (“Young Waits for Old”)

If *requesting* txn has higher priority than *holding* txn

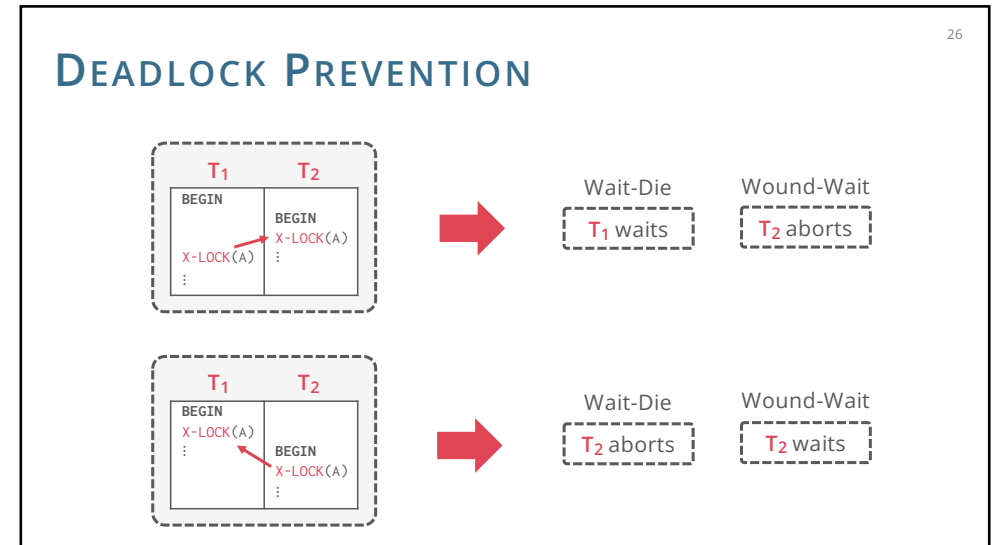
Then *holding* txn **aborts** and releases locks

Else *requesting* txn **waits**

24



25



26

DEADLOCK PREVENTION

27

*Why do these schemes guarantee no deadlocks?*

Only one "type" of direction allowed when waiting for a lock

*When a transaction restarts, what is its (new) priority?*

Its original timestamp. Why?

27

SUMMARY

28

**ACID Transactions**

- Atomicity:** All or nothing
- Consistency:** Only valid data
- Isolation:** No interference
- Durability:** Committed data persists

**Concurrency Control**

- Prevent anomalous schedules
- Locks + protocol (2PL, Strict 2PL) guarantees conflict serializability
- Deadlock detection and deadlock prevention

**Serializability**

- Serializable schedules
- Conflict & view serializability
- Checking for conflict serializability

28