

Conjunctive Queries: Evaluation and Static Analysis

(Chapter 14 and 15 of DBT)

[DBT] Database Theory, <https://github.com/pdm-book/community>

Semantics of Conjunctive Queries

- A **match** of a conjunctive query $Q(x_1, \dots, x_k) :- \text{body}$ in a database D is a homomorphism h from the set of atoms **body** to the set of atoms D
- The **answer** to $Q(x_1, \dots, x_k) :- \text{body}$ over D is the set of k -tuples

$$Q(D) := \{h(x_1), \dots, h(x_k) \mid h \text{ is a match of } Q \text{ in } D\}$$
- The answer consists of the witnesses for the **distinguished variables** of Q

Pattern Matching Problem

List the airlines that fly directly from London to Glasgow

Flight(VIE, LHR, BA),	Airport(VIE, Vienna),
Flight(LHR, EDI, BA),	Airport(LHR, London),
Flight(LGW, GLA, U2),	Airport(LGW, London),
Flight(LCA, VIE, OS),	Airport(LCA, Larnaca),
	Airport(GLA, Glasgow),
	Airport(EDI, Edinburgh)

$Q(z) :- \text{Airport}(x, \text{London}), \text{Airport}(y, \text{Glasgow}), \text{Flight}(x, y, z)$

Pattern Matching Problem

List the airlines that fly directly from London to Glasgow

Flight(VIE, LHR, BA),	Airport(VIE, Vienna),
Flight(LHR, EDI, BA),	Airport(LHR, London),
Flight(LGW, GLA, U2),	Airport(LGW, London),
Flight(LCA, VIE, OS),	Airport(LCA, Larnaca),
	Airport(GLA, Glasgow),
	Airport(EDI, Edinburgh)

$\{x \mapsto \text{LGW}, y \mapsto \text{GLA}, z \mapsto \text{U2},$
 $\text{London} \mapsto \text{London}, \text{Glasgow} \mapsto \text{Glasgow}\}$

$Q(z) :- \text{Airport}(x, \text{London}), \text{Airport}(y, \text{Glasgow}), \text{Flight}(x, y, z)$

Query Evaluation

- Understand the complexity of evaluating a conjunctive query over a database
- What to measure? Queries may have a large output, and it would be misleading to count the output as “complexity”
- We therefore consider the following decision problem for **CQ**

CQ-Evaluation

Input: a database D , a CQ $Q(x_1, \dots, x_k) :- \text{body}$, and a tuple (a_1, \dots, a_k) of values

Question: $(a_1, \dots, a_k) \in Q(D)$?

combined complexity

Data Complexity of Query Evaluation

- Measures the complexity in terms of the size of the database - the query is fixed
- Meaningful in practice since the database is usually much bigger than the query
- We consider the following decision problem for a fixed CQ $Q(x_1, \dots, x_k) :- \text{body}$

Q-Evaluation

Input: a database D , and a tuple (a_1, \dots, a_k) of values

Question: $(a_1, \dots, a_k) \in Q(D)$?

Complexity of Query Evaluation

Theorem: CQ-Evaluation is NP-complete, and in PTIME in data complexity

Proof:

(NP-membership) Guess-and-check:

- Consider a database D , a CQ $Q(x_1, \dots, x_k) :- \text{body}$, and a tuple (a_1, \dots, a_k) of values
- Guess a substitution $h : \text{terms}(\text{body}) \rightarrow \text{terms}(D)$
- Verify that h is a match of Q in D , i.e., $h(\text{body}) \subseteq D$ and $(h(x_1), \dots, h(x_k)) = (a_1, \dots, a_k)$

(NP-hardness) Reduction from 3-colorability

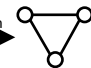
NP-hardness

(NP-hardness) Reduction from 3-colorability

3COL

Input: an undirected graph $G = (V, E)$

Question: is there a function $c : V \rightarrow \{R, G, B\}$ such that $(v, u) \in E \Rightarrow c(v) \neq c(u)$?

Lemma: G is 3-colorable iff G can be mapped to K_3 , i.e., $G \xrightarrow{\text{hom}}$ 

therefore, G is 3-colorable iff there is a match of Q_G in $D = \{E(x, y), E(y, z), E(z, x)\}$

the Boolean CQ that represents G

Complexity of Query Evaluation

Theorem: CQ-Evaluation is NP-complete, and in PTIME in data complexity

Proof:

(NP-membership) Guess-and-check:

- Consider a database D , a CQ $Q(x_1, \dots, x_k) :- \text{body}$, and a tuple (a_1, \dots, a_k) of values
- Guess a substitution $h : \text{terms}(\text{body}) \rightarrow \text{terms}(D)$
- Verify that h is a match of Q in D , i.e., $h(\text{body}) \subseteq D$ and $(h(x_1), \dots, h(x_k)) = (a_1, \dots, a_k)$

(NP-hardness) Reduction from 3-colorability

(in PTIME) For every substitution $h : \text{terms}(\text{body}) \rightarrow \text{terms}(D)$, check if $h(\text{body}) \subseteq D$ and $(h(x_1), \dots, h(x_k)) = (a_1, \dots, a_k)$

Static Analysis

CQ-Satisfiability

Input: a conjunctive query Q

Question: is there a database D such that $Q(D)$ is non-empty?

- If the answer is no, then the input query Q makes no sense
- CQ-Evaluation becomes trivial - the answer is always NO!

Static Analysis

CQ-Equivalence

Input: two conjunctive queries Q_1 and Q_2

Question: $Q_1 \equiv Q_2$? or $Q_1(D) = Q_2(D)$ for every database D ?

- Replace a query Q_1 with a query Q_2 that is easier to evaluate
- But, we have to be sure that $Q_1(D) = Q_2(D)$ for every database D

Static Analysis

CQ-Containment

Input: two conjunctive queries Q_1 and Q_2

Question: $Q_1 \subseteq Q_2$? or $Q_1(D) \subseteq Q_2(D)$ for every database D ?

- Equivalence boils down to two containment checks
- Clearly, $Q_1(D) = Q_2(D)$ iff $Q_1(D) \subseteq Q_2(D)$ and $Q_2(D) \subseteq Q_1(D)$

Complexity of Static Analysis

CQ-Satisfiability

Input: a conjunctive query Q

Question: is there a database D such that $Q(D)$ is non-empty?

CQ-Equivalence

Input: two conjunctive queries Q_1 and Q_2

Question: $Q_1 \equiv Q_2$? or $Q_1(D) = Q_2(D)$ for every database D ?

CQ-Containment

Input: two conjunctive queries Q_1 and Q_2

Question: $Q_1 \subseteq Q_2$? or $Q_1(D) \subseteq Q_2(D)$ for every database D ?

Canonical Database

- Convert a conjunctive query Q into a database $D[Q]$ - the **canonical database** of Q
- Given a conjunctive query of the form $Q(x) :- \text{body}$, $D[Q]$ is obtained from body by replacing each variable x with a new value $c(x) = x_$
- E.g., given $Q(x,y) :- R(x,y), P(y,z,w), R(z,x)$, then $D[Q] = \{R(x_1,y_1), P(y_1,z_1,w_1), R(z_1,x_1)\}$
- Note:** The mapping $c : \{\text{variables in body}\} \rightarrow \{\text{new values}\}$ is a **bijection**, where $c(\text{body}) = D[Q]$ and $c^{-1}(D[Q]) = \text{body}$

Satisfiability of CQs

CQ-Satisfiability

Input: a conjunctive query Q

Question: is there a database D such that $Q(D)$ is non-empty?

Theorem: A conjunctive query Q is always satisfiable

Proof: Due to its canonical database - $Q(D[Q])$ is trivially non-empty

Equivalence and Containment of CQs

CQ-Equivalence

Input: two conjunctive queries Q_1 and Q_2

Question: $Q_1 \equiv Q_2$? or $Q_1(D) = Q_2(D)$ for every database D ?

CQ-Containment

Input: two conjunctive queries Q_1 and Q_2

Question: $Q_1 \subseteq Q_2$? or $Q_1(D) \subseteq Q_2(D)$ for every database D ?

$$Q_1 \equiv Q_2 \text{ iff } Q_1 \subseteq Q_2 \text{ and } Q_2 \subseteq Q_1$$

$$Q_1 \subseteq Q_2 \text{ iff } Q_1 \equiv (Q_1 \wedge Q_2)$$

...thus, we can safely focus on **CQ-Containment**

Homomorphism Theorem

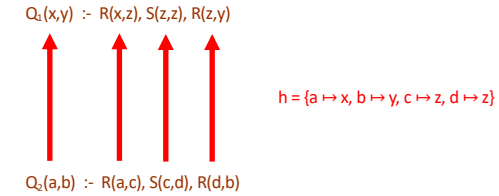
A **query homomorphism** from $Q_1(x_1, \dots, x_k) \text{ :- } \text{body}_1$ to $Q_2(y_1, \dots, y_k) \text{ :- } \text{body}_2$ is a substitution $h : \text{terms}(\text{body}_1) \rightarrow \text{terms}(\text{body}_2)$ such that:

1. h is a homomorphism from body_1 to body_2
2. $(h(x_1), \dots, h(x_k)) = (y_1, \dots, y_k)$

Homomorphism Theorem: Let Q_1 and Q_2 be conjunctive queries. It holds that:

$Q_1 \subseteq Q_2$ iff there exists a query homomorphism from Q_1 to Q_2

Homomorphism Theorem: Example



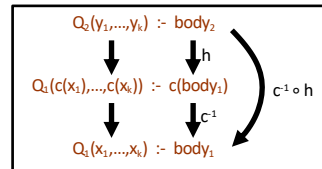
- h is a query homomorphism from Q_2 to $Q_1 \Rightarrow Q_1 \subseteq Q_2$
- But, there is no homomorphism from Q_1 to $Q_2 \Rightarrow Q_1 \subsetneq Q_2$

Homomorphism Theorem: Proof

Assume that $Q_1(x_1, \dots, x_k) \text{ :- } \text{body}_1$ and $Q_2(y_1, \dots, y_k) \text{ :- } \text{body}_2$

$(\Rightarrow) Q_1 \subseteq Q_2 \Rightarrow$ there exists a query homomorphism from Q_2 to Q_1

- Clearly, $(c(x_1), \dots, c(x_k)) \in Q_1(D[Q_1])$ - recall that $D[Q_1] = c(\text{body}_1)$
- Since $Q_1 \subseteq Q_2$, we conclude that $(c(x_1), \dots, c(x_k)) \in Q_2(D[Q_1])$
- Therefore, there exists a homomorphism h such that $h(\text{body}_2) \subseteq D[Q_1] = c(\text{body}_1)$ and $h((y_1, \dots, y_k)) = (c(x_1), \dots, c(x_k))$
- By construction, $c^{-1}(c(\text{body}_1)) = \text{body}_1$ and $c^{-1}((c(x_1), \dots, c(x_k))) = (x_1, \dots, x_k)$
- Therefore, $c^{-1} \circ h$ is a query homomorphism from Q_2 to Q_1

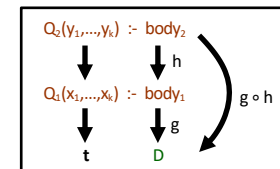


Homomorphism Theorem: Proof

Assume that $Q_1(x_1, \dots, x_k) \text{ :- } \text{body}_1$ and $Q_2(y_1, \dots, y_k) \text{ :- } \text{body}_2$

$(\Leftarrow) Q_1 \subseteq Q_2 \Leftarrow$ there exists a query homomorphism from Q_2 to Q_1

- Consider a database D , and a tuple t such that $t \in Q_1(D)$
- We need to show that $t \in Q_2(D)$
- Clearly, there exists a homomorphism g such that $g(\text{body}_1) \subseteq D$ and $g((x_1, \dots, x_k)) = t$
- By hypothesis, there exists a query homomorphism h from Q_2 to Q_1
- Therefore, $g(h(\text{body}_2)) \subseteq D$ and $g(h((y_1, \dots, y_k))) = t$, which implies that $t \in Q_2(D)$



Existence of a Query Homomorphism

Theorem: Let Q_i and Q_j be conjunctive queries. The problem of deciding whether there exists a query homomorphism from Q_j to Q_i is NP-complete

Proof:

(NP-membership) Guess a substitution, and verify that is a query homomorphism

(NP-hardness) Easy reduction from CQ-Evaluation

By applying the homomorphism theorem we get that:

Corollary: CQ-Equivalence and CQ-Containment are NP-complete