

UNIVERSITY OF EDINBURGH  
SCHOOL OF INFORMATICS  
INFR11199 - ADVANCED DATABASE SYSTEMS (SPRING 2024)

Tutorial Sheet 3

1. (Sorting and Hashing) Suppose the size of a page is 4 KB, and the size of the memory buffer is 1 MB (1024 KB).
  - (a) We have a relation of size 800 KB. How many page I/Os are required to sort this relation and write the sorted relation back to disk?
  - (b) We have a relation of size 5000 KB. How many page I/Os are required to sort this relation and write the sorted relation back to disk?
  - (c) What is the size of the largest relation that would need two passes to sort?
  - (d) What is the size of the largest relation we can possibly hash in two passes (i.e., with just one partitioning phase)?
  - (e) Suppose we have a relation of size 3000 KB. We are executing a `DISTINCT` query on a column `age`, which has only two distinct values, evenly distributed. Would sorting or hashing be better here, and why?
  - (f) Now suppose we were executing a `GROUP BY` on `age` instead. Would sorting or hashing be better here, and why?
2. (Joins) Consider the following database of students and assignment submissions and the SQL query:

```
CREATE TABLE Students (  
    student_id INTEGER PRIMARY KEY,  
    ...  
);  
CREATE TABLE Assignments(  
    assignment_number INTEGER,  
    student_id INTEGER REFERENCES Students(student_id),  
    ...  
);  
SELECT *  
FROM Students, Assignments  
WHERE Students.student_id = Assignments.student_id;
```

Assume the following:

- **Students** has 20 pages, with 200 records per page
- **Assignments** has 40 pages, with 250 records per page.

- (a) What is the I/O cost of a simple nested loop join for **Students** ⋈ **Assignments**?
- (b) What is the I/O cost of a simple nested loop join for **Assignments** ⋈ **Students**?
- (c) What is the I/O cost of a block nested loop join for **Students** ⋈ **Assignments**?  
Assume our buffer size is  $B = 12$  pages.
- (d) What is the I/O cost of a block nested loop join for **Assignments** ⋈ **Students**?  
Assume our buffer size is  $B = 12$  pages.
- (e) What is the I/O cost of an Index-Nested Loop Join for **Students** on ⋈ **Assignments**?

Assume we have a *clustered* variant B index on **Assignments.student\_id**, in the form of a height 2 B+ tree. Assume that: index (non-leaf) nodes and leaf pages are not cached; all hits are on the same leaf page; and all hits are also on the same data page. only one matching page

- (f) Now assume we have an *unclustered* variant B index on **Assignments.student\_id**, in the form of a height 2 B+ tree. Assume that index node pages and leaf pages are never cached, and we only need to read the relevant leaf page once for each record of **Students**, and all hits are on the same leaf page.  
What is the I/O cost of an Index-Nested Loop Join for **Students** ⋈ **Assignments**?  
Hint: The foreign key in **Assignments** may play a role in how many accesses we do per record. no info on total number of matching tuples
- (g) What is the cost of an *unoptimized* sort-merge join for **Students** ⋈ **Assignments**?  
Assume we have  $B = 12$  buffer pages.
- (h) What is the cost of an *optimized* sort-merge join for **Students** ⋈ **Assignments**?  
Assume we have  $B = 12$  buffer pages.
- (i) In the previous question, we had a buffer of  $B = 12$  pages. If we shrank  $B$  enough, the answer we got might change.  
How small can the buffer  $B$  be without changing the I/O cost answer we got?
- (j) What is the I/O cost of Grace Hash Join on these tables?  
Assume uniform hash partitioning and a buffer pool consisting of  $B = 6$  pages.