

UNIVERSITY OF EDINBURGH
SCHOOL OF INFORMATICS
INFR11199 - ADVANCED DATABASE SYSTEMS (SPRING 2024)

Tutorial Sheet 3

1. (Sorting and Hashing) Suppose the size of a page is 4 KB, and the size of the memory buffer is 1 MB (1024 KB).

- (a) We have a relation of size 800 KB. How many page I/Os are required to sort this relation and write the sorted relation back to disk?

Solution: 400 ($= 200 + 200$).

200 to read in, 200 to write out. Since the relation is small enough to completely fit into the buffer, we only need to read it in, sort it (no I/Os required for sorting), then write the sorted pages back to disk.

- (b) We have a relation of size 5000 KB. How many page I/Os are required to sort this relation and write the sorted relation back to disk?

Solution: 5000 ($= 2 * 1250 * \text{number of passes}$).

2 passes. 5000 KB with 4KB per page means 1250 pages are needed to store the relation. We have $B = 1024 / 4 = 256$ pages in our buffer.

Number of Passes $= 1 + \lceil \log_{256} \lceil 1250/256 \rceil \rceil = 2$.

- (c) What is the size of the largest relation that would need two passes to sort?

Solution: 261,120 KB. ($255 * 256$ pages).

- (d) What is the size of the largest relation we can possibly hash in two passes (i.e., with just one partitioning phase)?

Solution: 261,120 KB.

- (e) Suppose we have a relation of size 3000 KB. We are executing a **DISTINCT** query on a column **age**, which has only two distinct values, evenly distributed. Would sorting or hashing be better here, and why?

Solution: Hashing, which allows us to remove duplicates early on and potentially improve performance (in this case, we might be able to finish in 1 pass, instead of 2 for sorting).

- (f) Now suppose we were executing a **GROUP BY** on **age** instead. Would sorting or hashing be better here, and why?

Solution: Sorting because hashing won't work; each partition is larger than memory, so no amount of hash partitioning will suffice.

2. (Joins) Consider the following database of students and assignment submissions and the SQL query:

```
CREATE TABLE Students (  
    student_id INTEGER PRIMARY KEY,  
    ...  
);  
CREATE TABLE Assignments(  
    assignment_number INTEGER,  
    student_id INTEGER REFERENCES Students(student_id),  
    ...  
);  
SELECT *  
FROM Students, Assignments  
WHERE Students.student_id = Assignments.student_id;
```

Assume the following:

- **Students** has 20 pages, with 200 records per page
- **Assignments** has 40 pages, with 250 records per page.

- (a) What is the I/O cost of a simple nested loop join for **Students** ⋈ **Assignments**?

Solution: 160,020 I/Os.
The cost of a SNLJ is: $\#pages(S) + \#records(S) \cdot \#pages(A)$.
Plugging in the numbers gives us $20 + (20 \cdot 200) \cdot 40 = 160,020$ I/Os.

- (b) What is the I/O cost of a simple nested loop join for **Assignments** ⋈ **Students**?

Solution: 200,040 I/Os.

The cost of a SNLJ is: $\#pages(A) + \#records(A) \cdot \#pages(S)$.

Plugging in the numbers gives us $40 + (40 \cdot 250) \cdot 20 = 200,040$ I/Os.

- (c) What is the I/O cost of a block nested loop join for **Students** \bowtie **Assignments**? Assume our buffer size is $B = 12$ pages.

Solution: 100 I/Os.

Since **Students** is the outer table, we calculate the number of blocks of **Students**: $\#pages(S) / (B-2) = 20 / 10 = 2$. Thus, the final cost is $\#pages(S)$ plus 2 passes through all pages(A), or $20 + 2 \cdot 40 = 100$ I/Os.

- (d) What is the I/O cost of a block nested loop join for **Assignments** \bowtie **Students**? Assume our buffer size is $B = 12$ pages.

Solution: 120 I/Os.

Since **Assignments** is the outer table, we calculate the number of blocks of **Assignments**: $\#pages(A) / (B-2) = 40 / 10 = 4$. Thus, the final cost is $\#pages(A)$ plus 4 passes through all pages(S), or $40 + 4 \cdot 20 = 120$ I/Os.

- (e) What is the I/O cost of an Index-Nested Loop Join for **Students** on \bowtie **Assignments**?

Assume we have a *clustered* variant B index on **Assignments.student_id**, in the form of a height 2 B+ tree. Assume that: index (non-leaf) nodes and leaf pages are not cached; all hits are on the same leaf page; and all hits are also on the same data page.

Solution: 16,020 I/Os.

The formula is $\#pages(S) + \#records(S) \cdot (\text{cost of index lookup})$.

The cost of index lookup is 3 I/Os to access the leaf, and 1 I/O to access the data page for all matching records.

So the total cost is $20 + 4000 \cdot 4 = 16,020$ I/Os.

- (f) Now assume we have an *unclustered* variant B index on **Assignments.student_id**, in the form of a height 2 B+ tree. Assume that index node pages and leaf pages are never cached, and we only need to read the relevant leaf page once for each record of **Students**, and all hits are on the same leaf page.

What is the I/O cost of an Index-Nested Loop Join for **Students** \bowtie **Assignments**?

Hint: The foreign key in **Assignments** may play a role in how many accesses we do per record.

Solution: 22,020 I/Os.

The formula is $\#pages(S) + \#records(S) * (\text{cost of index lookup})$.

This time though, the cost of index lookup is 3 I/Os to access the leaf, and 1 I/O to access the data page for *each matching record*.

How many records match per key? We actually haven't told you! But, we do know that we will eventually have to access each record exactly once (since each **Assignments** is foreign-keyed on a **student_id** – so there will be $\#records(A) = 10,000$ data page lookups, one for each row).

So the total cost is $20 + 4000 \cdot 3 + 10000 = 22,020$ I/Os.

- (g) What is the cost of an *unoptimized* sort-merge join for **Students** \bowtie **Assignments**? Assume we have $B = 12$ buffer pages.

Solution: 300 I/Os.

The formula is $(\text{cost of sorting } S) + (\text{cost of sorting } A) + \#pages(S) + \#pages(A)$.

For sorting S : The first pass will make two runs, which is mergeable in one merge pass; thus, we need two passes.

For sorting A : The first pass will make four runs, which is mergeable in one merge pass; thus, we need two passes.

Thus the total cost is $(2 \cdot 2\#pages(S)) + (2 \cdot 2\#pages(A)) + \#pages(S) + \#pages(A) = 5(\#pages(S) + \#pages(A)) = 5 \cdot 60 = 300$ I/Os.

- (h) What is the cost of an *optimized* sort-merge join for **Students** \bowtie **Assignments**? Assume we have $B = 12$ buffer pages.

Solution: 180 I/Os.

The difference from the above question is that we will skip the last write in the external sorting phase, and the initial read in the sort-merge phase. For this to be possible, all the runs of S and A in the last phase of external sorting should be able to fit into memory together. From the previous question, we know there are $2 + 4 = 6$ runs, which fits just fine in our buffer of 12 pages.

The total cost is $300 - 2\#pages(S) - 2\#pages(A) = 300 - 120 = 180$ I/Os.

- (i) In the previous question, we had a buffer of $B = 12$ pages. If we shrank B enough, the answer we got might change. How small can the buffer B be without changing the I/O cost answer we got?

Solution: The restriction for optimized sort-merge join is that the number of final runs of S and A can both fit in memory simultaneously (i.e., the

number of runs of S + the number of runs of $A \leq B - 1$). We had $2 + 4$ runs last time, which fit comfortably in $12 - 1$ buffer pages (recall that one page is reserved for output).

What about $B = 11$? We would still have $2 + 4 \leq 11 - 1$ runs.

What about $B = 10$? We would still have $2 + 4 \leq 10 - 1$ runs.

What about $B = 9$? Now we have 3 runs for S and 5 runs for A , which just exactly fits in $9 - 1$ buffer pages.

Since 9 buffer pages fits perfectly, any smaller would force more merge passes and thus more I/Os.

- (j) What is the I/O cost of Grace Hash Join on these tables?

Assume uniform hash partitioning and a buffer pool consisting of $B = 6$ pages.

Solution: 180 I/Os.

For Grace Hash Join, we have to walk through what the partition sizes are like for each phase, one phase at a time. In the partitioning phase, we will proceed as in external hashing. We will load in 1 page at a time and hash it into $B - 1 = 5$ partitions. This means the 20 pages of S get split into 4 pages per partition, and the 40 pages of A get split into 8 pages per partition.

Do we need to recursively partition? No! Remember that the stopping condition is that any table's partition fits in $B - 2 = 4$ buffer pages; the partitions of S satisfy this.

In the hash joining phase, the I/O cost is simply the total number of pages across all partitions – we read all of these in exactly once.

Thus the final I/O cost is $20 + 20$ for partitioning S , $40 + 40$ for partitioning A , and $20 + 40$ for the hash join, for a total cost of 180 I/Os.