

UNIVERSITY OF EDINBURGH
SCHOOL OF INFORMATICS
INFR11199 - ADVANCED DATABASE SYSTEMS (SPRING 2024)

Tutorial Sheet 4 - Playing with PostgreSQL

The purpose of this practical sheet is to familiarise you with the query execution engine of PostgreSQL. In particular, you will analyse a few queries and answer questions regarding their performance when turning different knobs of the execution engine. To answer the questions, you might find the following documentation links useful:

- Documentation of `EXPLAIN ANALYZE`:
<https://www.postgresql.org/docs/14/sql-explain.html>.
- Making sense of the `EXPLAIN ANALYZE` output:
<https://www.postgresql.org/docs/14/performance-tips.html>.
- PostgreSQL query planner documentation:
<https://www.postgresql.org/docs/14/runtime-config-query.html>.
- How to create an index:
<https://www.postgresql.org/docs/14/sql-createindex.html>.
- The system table `pg_class`:
<https://www.postgresql.org/docs/current/catalog-pg-class.html>.

Prerequisites:

- Install PostgreSQL on your machine and start a PostgreSQL server (plenty of instructions online on how to do this, e.g., <http://postgresguide.com/setup/install.html>; any version will work). Make sure the command-line tool `psql` is working and you can use it to create tables and run queries.
- Download the bay-area-bike-sharing dataset from the course webpage. Unzip the archive and import the data into PostgreSQL using the provided scripts (e.g., by typing the command `psql < import.sql`).

1. EXPLAIN and ANALYZE

For the following questions consider the query below:

```
SELECT * FROM trip WHERE bike_id = 10;
```

- (a) Provide the PostgreSQL execution plan of the query and the SQL statement you use to generate the result.



Based on the execution plan:

- i. What was the estimated cost (in arbitrary units)? _____
- ii. What was the total runtime (in ms)? _____
- iii. What was the estimated number of tuples to be output? _____
- iv. What was the actual number of output tuples? _____

- (b) Create an index on the attribute `bike_id` on the table `trip`. Provide the SQL statement for that and the new execution plan of the query.

Based on the execution plan:

- i. What was the estimated cost (in arbitrary units)? Faster than without index _____
- ii. What was the total runtime (in ms)? _____

- (c) Use the table `pg_class` to answer the following questions.

- i. How many pages are used to store the index you created on table `trip`? Provide the answer and the query you use to generate the answer.

- ii. How many tuples are in the index you created on column `bike_id`? Provide the answer and the query you use to generate the answer.

- iii. How many tuples are in the table `weather`, according to `pg_class`?

- iv. In the table `weather`, delete all records of which date is earlier than '2013-10-01'. Provide the SQL statement you use.

- v. After deletion, rerun your query from step 3. Is the new result equal to the result of running `SELECT COUNT(*) FROM weather`?

☐ Yes ☐ No

- vi. ANALYZE is a Postgres function used to collect statistics about a database. You want to use it especially after considerable number of modifications happen to that database. Run ANALYZE, and then rerun your query from step 3 again. Is the new result equal to the result of running `SELECT COUNT(*) FROM weather`?

☐ Yes ☐ No

2. Using Indexes

In this question, we will learn the conditions under which indexes may or may not be used by the query optimizer.

- (a) Create an index on the column `start_station_name` on the table `trip`. Provide the SQL command you use.

- (b) For each of those queries, answer Yes if the index you created on `trip.start_station_name` was used in the execution plan, or No otherwise:

- i. `SELECT * FROM trip
WHERE start_station_name like 'San';`

☐ Yes ☐ No

- ii. `SELECT * FROM trip
WHERE start_station_name like '%San';`

☐ Yes ☐ No

- iii. `SELECT * FROM trip
WHERE start_station_name BETWEEN 'San_Francisco' AND 'San_Jose'
AND end_station_name > 'San';`

☐ Yes ☐ No

iv. `SELECT * FROM trip`
`WHERE start_station_name BETWEEN 'San Francisco' AND 'San Jose'`
`OR end_station_name > 'San';`

☐ Yes ☐ No

(c) Make sure you still have an index on the column `trip.bike_id` (you can verify this using `\di` in `psql`). For each of those queries, answer which indexes are used in their execution plans.

i. `SELECT * FROM trip`
`WHERE start_station_name BETWEEN 'San Francisco' AND 'San Jose'`
`AND bike_id < 10;`

- ☐ Only the index on `start_station_name` was used.
- ☐ Only the index on `bike_id` was used.
- ☐ Both indexes were used.
- ☐ None of the indexes were used.

ii. `SELECT * FROM trip`
`WHERE start_station_name BETWEEN 'San Francisco' AND 'San Jose'`
`AND bike_id < 500;`

- ☐ Only the index on `start_station_name` was used.
- ☐ Only the index on `bike_id` was used.
- ☐ Both indexes were used.
- ☐ None of the indexes were used.

iii. `SELECT * FROM trip`
`WHERE start_station_name BETWEEN 'San Francisco' AND 'San Jose'`
`AND bike_id BETWEEN 500 AND 510;`

- ☐ Only the index on `start_station_name` was used.
- ☐ Only the index on `bike_id` was used.
- ☐ Both indexes were used.
- ☐ None of the indexes were used.

iv. `SELECT * FROM trip`
`WHERE start_station_name > 'San Francisco'`
`AND bike_id < 500;`

- ☐ Only the index on `start_station_name` was used.
- ☐ Only the index on `bike_id` was used.
- ☐ Both indexes were used.
- ☐ None of the indexes were used.

(d) Answer the questions below for the query:

```
SELECT * FROM trip
WHERE bike_id BETWEEN 10 AND 20;
```

i. Was the index on `bike_id` used? ☐ Yes ☐ No

- ii. What percentage of the total records in the table `trip` was returned?
Provide a percent and retain two significant figures. _____

(e) Answer the questions below for the query:

```
SELECT * FROM trip
WHERE bike_id > 10;
```

- i. Was the index on `bike_id` used? ☐ Yes ☐ No
- ii. What percentage of the total records in the table `trip` was returned?
Provide a percent and retain two significant figures. _____

(f) Answer the questions below for the query:

```
SELECT * FROM trip
WHERE bike_id > 10 ORDER BY start_time;
```

- i. Which method was used for sorting? _____
- ii. Where did the sorting happen? ☐ Memory ☐ Disk
- iii. How much space was used for sorting? _____
- iv. What was the total runtime (in ms)? _____

(g) Display PostgreSQL working memory with `SHOW work_mem;`. Increase PostgreSQL working memory with the command `SET work_mem = '128MB';`. For the same query from part vi., answer the following questions:

- i. Which method was used for sorting? _____
- ii. Where did the sorting happen? ☐ Memory ☐ Disk
- iii. How much space was used for sorting? _____
- iv. What was the total runtime (in ms)? _____

(h) Execute the command `RESET work_mem;` to get PostgreSQL working memory back to the default value (or your answers for the next questions will turn out wrong).

3. Joins

In this question, we will learn about different methods used by PostgreSQL for executing joins. Make sure you reset `work_mem` to its default value (i.e., `RESET work_mem;`).

Answer the following questions based on the query below:

```
SELECT trip.*, station.city
FROM trip, station
WHERE trip.start_station_id = station.station_id
AND bike_id < 200;
```

- (a) Provide the query plan for the above query.

Based on the execution plan:

- i. Which join method was used? _____
- ii. What was the estimated cost (in arbitrary units)? _____
- iii. What was the total runtime (in ms)? _____

- (b) Execute the command `SET enable_hashjoin = false;` to disable hash joins. Provide the new query plan.

Based on the execution plan:

- i. Which join method was used? _____

- ii. What was the estimated cost (in arbitrary units)? _____
- iii. What was the total runtime (in ms)? _____

- (c) Execute the command `SET enable_mergejoin = false;` to disable merge joins. Provide the new query plan.

Based on the execution plan:

- i. Which join method was used? _____
- ii. What was the estimated cost (in arbitrary units)? _____
- iii. What was the total runtime (in ms)? _____

- (d) Execute the command `SET enable_indexscan = false; SET enable_bitmapscan = false;` to disable index scans. Give the new plan.

Based on the execution plan:

- i. Which join method was used? _____
- ii. What was the estimated cost (in arbitrary units)? _____
- iii. What was the total runtime (in ms)? _____

(e) Execute these commands to re-enable the different joins.

```
RESET enable_mergejoin;  
RESET enable_hashjoin;  
RESET enable_indexscan;  
RESET enable_bitmapscan;
```