

UNIVERSITY OF EDINBURGH
SCHOOL OF INFORMATICS
INFR11199 - ADVANCED DATABASE SYSTEMS (SPRING 2024)

Tutorial Sheet 4 - Playing with PostgreSQL

The purpose of this practical sheet is to familiarise you with the query execution engine of PostgreSQL. In particular, you will analyse a few queries and answer questions regarding their performance when turning different knobs of the execution engine. To answer the questions, you might find the following documentation links useful:

- Documentation of `EXPLAIN ANALYZE`:
<https://www.postgresql.org/docs/14/sql-explain.html>.
- Making sense of the `EXPLAIN ANALYZE` output:
<https://www.postgresql.org/docs/14/performance-tips.html>.
- PostgreSQL query planner documentation:
<https://www.postgresql.org/docs/14/runtime-config-query.html>.
- How to create an index:
<https://www.postgresql.org/docs/14/sql-createindex.html>.
- The system table `pg_class`:
<https://www.postgresql.org/docs/current/catalog-pg-class.html>.

Prerequisites:

- Install PostgreSQL on your machine and start a PostgreSQL server (plenty of instructions online on how to do this, e.g., <http://postgresguide.com/setup/install.html>; any version will work). Make sure the command-line tool `psql` is working and you can use it to create tables and run queries.
- Download the bay-area-bike-sharing dataset from the course webpage. Unzip the archive and import the data into PostgreSQL using the provided scripts (e.g., by typing the command `psql < import.sql`).

1. EXPLAIN and ANALYZE

For the following questions consider the query below:

```
SELECT * FROM trip WHERE bike_id = 10;
```

- (a) Provide the PostgreSQL execution plan of the query and the SQL statement you use to generate the result.

Solution:

```
EXPLAIN ANALYZE SELECT * FROM trip WHERE bike_id = 10;

QUERY PLAN
-----
Gather  (cost=1000.00..14191.37 rows=790 width=80) (actual time
=2.556..143.164 rows=248 loops=1)
  Workers Planned: 2
  Workers Launched: 2
  -> Parallel Seq Scan on trip  (cost=0.00..13112.37 rows=329 width=80)
      (actual time=4.328..127.060 rows=83 loops=3)
        Filter: (bike_id = 10)
        Rows Removed by Filter: 223237
Planning Time: 0.824 ms
Execution Time: 143.286 ms
```

Based on the execution plan:

- | | |
|------------------------------------------------------------|-----------------|
| i. What was the estimated cost (in arbitrary units)? | <u>14191.87</u> |
| ii. What was the total runtime (in ms)? | <u>143.286</u> |
| iii. What was the estimated number of tuples to be output? | <u>790</u> |
| iv. What was the actual number of output tuples? | <u>248</u> |

- (b) Create an index on the attribute `bike_id` on the table `trip`. Provide the SQL statement for that and the new execution plan of the query.

Solution:

```
CREATE INDEX idx_bike_id ON trip(bike_id);
```

```
QUERY PLAN
```

```
-----  
Bitmap Heap Scan on trip  (cost=18.55..2424.94 rows=790 width=80) (actual  
  time=0.432..1.039 rows=248 loops=1)  
  Recheck Cond: (bike_id = 10)  
  Heap Blocks: exact=234  
-> Bitmap Index Scan on idx_bike_id  (cost=0.00..18.35 rows=790 width  
    =0) (actual time=0.335..0.335 rows=248 loops=1)  
      Index Cond: (bike_id = 10)  
Planning Time: 6.332 ms  
Execution Time: 1.154 ms
```

Based on the execution plan:

- | | |
|------------------------------------------------------|----------------|
| i. What was the estimated cost (in arbitrary units)? | <u>2424.94</u> |
| ii. What was the total runtime (in ms)? | <u>1.154</u> |

- (c) Use the table `pg_class` to answer the following questions.

- i. How many pages are used to store the index you created on table `trip`? Provide the answer and the query you use to generate the answer.

Solution:

```
SELECT relpages FROM pg_class  
WHERE relname = 'idx_bike_id';
```

```
relpages  
-----  
      1840
```

- ii. How many tuples are in the index you created on column `bike_id`? Provide the answer and the query you use to generate the answer.

Solution:

```
SELECT reltuples FROM pg_class  
WHERE relname = 'idx_bike_id';
```

```
reltuples  
-----  
    669959
```

- iii. How many tuples are in the table `weather`, according to `pg_class`?

Solution:

```
SELECT reltuples FROM pg_class
WHERE relname = 'weather';
```

```
    reltuples
-----
         3665
```

- iv. In the table `weather`, delete all records of which date is earlier than '2013-10-01'. Provide the SQL statement you use.

Solution:

```
DELETE FROM weather WHERE date < '2013-10-01';
```

- v. After deletion, rerun your query from step 3. Is the new result equal to the result of running `SELECT COUNT(*) FROM weather`?

☐ Yes ☒ No

- vi. `ANALYZE` is a Postgres function used to collect statistics about a database. You want to use it especially after considerable number of modifications happen to that database. Run `ANALYZE`, and then rerun your query from step 3 again. Is the new result equal to the result of running `SELECT COUNT(*) FROM weather`?

☒ Yes ☐ No

2. Using Indexes

In this question, we will learn the conditions under which indexes may or may not be used by the query optimizer.

- (a) Create an index on the column `start_station_name` on the table `trip`. Provide the SQL command you use.

Solution:

```
CREATE INDEX idx_start_sta_name ON trip ( start_station_name );
```

- (b) For each of those queries, answer Yes if the index you created on `trip.start_station_name` was used in the execution plan, or No otherwise:

- i.

```
SELECT * FROM trip
WHERE start_station_name like 'San';
```

☒ Yes ☐ No

ii. `SELECT * FROM trip
WHERE start_station_name like '%San';`

☐ Yes ☒ No

iii. `SELECT * FROM trip
WHERE start_station_name BETWEEN 'San_Francisco' AND 'San_Jose'
AND end_station_name > 'San';`

☒ Yes ☐ No

iv. `SELECT * FROM trip
WHERE start_station_name BETWEEN 'San_Francisco' AND 'San_Jose'
OR end_station_name > 'San';`

☐ Yes ☒ No

(c) Make sure you still have an index on the column `trip.bike_id` (you can verify this using `\di` in `psql`). For each of those queries, answer which indexes are used in their execution plans.

i. `SELECT * FROM trip
WHERE start_station_name BETWEEN 'San_Francisco' AND 'San_Jose'
AND bike_id < 10;`

- ☐ Only the index on `start_station_name` was used.
☒ **Only the index on `bike_id` was used.**
☐ Both indexes were used.
☐ None of the indexes were used.

ii. `SELECT * FROM trip
WHERE start_station_name BETWEEN 'San_Francisco' AND 'San_Jose'
AND bike_id < 500;`

- ☒ **Only the index on `start_station_name` was used.**
☐ Only the index on `bike_id` was used.
☐ Both indexes were used.
☐ None of the indexes were used.

iii. `SELECT * FROM trip
WHERE start_station_name BETWEEN 'San_Francisco' AND 'San_Jose'
AND bike_id BETWEEN 500 AND 510;`

- ☐ Only the index on `start_station_name` was used.
☐ Only the index on `bike_id` was used.
☒ **Both indexes were used.**
☐ None of the indexes were used.

iv. `SELECT * FROM trip
WHERE start_station_name > 'San_Francisco'
AND bike_id < 500;`

- ☐ Only the index on `start_station_name` was used.
- ☐ Only the index on `bike_id` was used.
- ☐ Both indexes were used.
- ☒ **None of the indexes were used.**

(d) Answer the questions below for the query:

```
SELECT * FROM trip
WHERE bike_id BETWEEN 10 AND 20;
```

- i. Was the index on `bike_id` used? ☒ **Yes** ☐ No
- ii. What percentage of the total records in the table `trip` was returned? Provide a percent and retain two significant figures. 0.54% (3594 out of 669959)

(e) Answer the questions below for the query:

```
SELECT * FROM trip
WHERE bike_id > 10;
```

- i. Was the index on `bike_id` used? ☐ Yes ☒ **No**
- ii. What percentage of the total records in the table `trip` was returned? Provide a percent and retain two significant figures. 99.92% (669460 out of 669959)

(f) Answer the questions below for the query:

```
SELECT * FROM trip
WHERE bike_id > 10 ORDER BY start_time;
```

- i. Which method was used for sorting? external merge sort
- ii. Where did the sorting happen? ☐ Memory ☒ **Disk**
- iii. How much space was used for sorting? 22080kB per worker
- iv. What was the total runtime (in ms)? 891.222 ms (any number)

(g) Display PostgreSQL working memory with `SHOW work_mem;`. Increase PostgreSQL working memory with the command `SET work_mem = '128MB';`. For the same query from part vi., answer the following questions:

- i. Which method was used for sorting? quick sort
- ii. Where did the sorting happen? ☒ **Memory** ☐ Disk
- iii. How much space was used for sorting? 115,909kB
- iv. What was the total runtime (in ms)? 454.680 ms (any number)

(h) Execute the command `RESET work_mem;` to get PostgreSQL working memory back to the default value (or your answers for the next questions will turn out wrong).

3. Joins

In this question, we will learn about different methods used by PostgreSQL for executing joins. Make sure you reset `work_mem` to its default value (i.e., `RESET work_mem;`).

Answer the following questions based on the query below:

```
SELECT trip.*, station.city
FROM trip, station
WHERE trip.start_station_id = station.station_id
AND bike_id < 200;
```

(a) Provide the query plan for the above query.

Solution:

```
QUERY PLAN
-----
Hash Join  (cost=1093.33..11604.43 rows=58107 width=92) (actual time
=29.274..85.848 rows=58161 loops=1)
  Hash Cond: (trip.start_station_id = station.station_id)
  -> Bitmap Heap Scan on trip  (cost=1090.75..11440.09 rows=58107 width
=80) (actual time=29.108..52.497 rows=58161 loops=1)
    Recheck Cond: (bike_id < 200)
    Heap Blocks: exact=9541
    -> Bitmap Index Scan on idx_bike_id  (cost=0.00..1076.23 rows
=58107 width=0) (actual time=25.510..25.510 rows=58161 loops=
=1)
        Index Cond: (bike_id < 200)
  -> Hash  (cost=1.70..1.70 rows=70 width=14) (actual time=0.126..0.126
rows=70 loops=1)
        Buckets: 1024  Batches: 1  Memory Usage: 12kB
        -> Seq Scan on station  (cost=0.00..1.70 rows=70 width=14) (
actual time=0.023..0.068 rows=70 loops=1)
Planning Time: 0.900 ms
Execution Time: 89.936 ms
```

Based on the execution plan:

- | | |
|-------------------------------------------------------|------------------|
| i. Which join method was used? | <u>Hash join</u> |
| ii. What was the estimated cost (in arbitrary units)? | <u>11604.43</u> |
| iii. What was the total runtime (in ms)? | <u>89.936</u> |

(b) Execute the command `SET enable_hashjoin = false;` to disable hash joins. Provide the new query plan.

Solution:

```
QUERY PLAN
-----
Merge Join  (cost=18625.07..19497.02 rows=58107 width=92) (actual time
=83.404..124.602 rows=58161 loops=1)
  Merge Cond: (trip.start_station_id = station.station_id)
  -> Sort  (cost=18621.22..18766.49 rows=58107 width=80) (actual time
=83.322..99.226 rows=58161 loops=1)
```

```

Sort Key: trip.start_station_id
Sort Method: external merge  Disk: 5400kB
-> Bitmap Heap Scan on trip  (cost=1090.75..11440.09 rows=58107
width=80) (actual time=15.734..39.091 rows=58161 loops=1)
    Recheck Cond: (bike_id < 200)
    Heap Blocks: exact=9541
    -> Bitmap Index Scan on idx_bike_id  (cost=0.00..1076.23
rows=58107 width=0) (actual time=13.722..13.722 rows
=58161 loops=1)
        Index Cond: (bike_id < 200)
-> Sort  (cost=3.85..4.02 rows=70 width=14) (actual time=0.075..0.084
rows=70 loops=1)
    Sort Key: station.station_id
    Sort Method: quicksort  Memory: 28kB
    -> Seq Scan on station  (cost=0.00..1.70 rows=70 width=14) (
actual time=0.016..0.033 rows=70 loops=1)
Planning Time: 0.642 ms
Execution Time: 134.899 ms

```

Based on the execution plan:

- | | |
|-------------------------------------------------------|------------------------|
| i. Which join method was used? | <u>Sort-merge join</u> |
| ii. What was the estimated cost (in arbitrary units)? | <u>19497.02</u> |
| iii. What was the total runtime (in ms)? | <u>134.899</u> |

- (c) Execute the command `SET enable_mergejoin = false;` to disable merge joins. Provide the new query plan.

Solution:

QUERY PLAN

```

-----
Nested Loop  (cost=1090.90..20735.16 rows=58107 width=92) (actual time
=19.123..179.892 rows=58161 loops=1)
-> Bitmap Heap Scan on trip  (cost=1090.75..11440.09 rows=58107 width
=80) (actual time=19.095..43.152 rows=58161 loops=1)
    Recheck Cond: (bike_id < 200)
    Heap Blocks: exact=9541
    -> Bitmap Index Scan on idx_bike_id  (cost=0.00..1076.23 rows
=58107 width=0) (actual time=16.637..16.638 rows=58161 loops
=1)
        Index Cond: (bike_id < 200)
-> Index Scan using station_pkey on station  (cost=0.14..0.16 rows=1
width=14) (actual time=0.002..0.002 rows=1 loops=58161)
    Index Cond: (station_id = trip.start_station_id)
Planning Time: 0.384 ms
Execution Time: 185.202 ms

```

Based on the execution plan:

- | | |
|-------------------------------------------------------|-----------------------------------------|
| i. Which join method was used? | <u>Nested loops join w/ index scans</u> |
| ii. What was the estimated cost (in arbitrary units)? | <u>20735.16</u> |
| iii. What was the total runtime (in ms)? | <u>185.202</u> |

- (d) Execute the command `SET enable_indexscan = false;` `SET enable_bitmapscan = false;` to disable index scans. Give the new plan.

Solution:

QUERY PLAN

```
-----  
Nested Loop (cost=0.00..78164.76 rows=58107 width=92) (actual time  
=0.101..595.171 rows=58161 loops=1)  
  Join Filter: (trip.start_station_id = station.station_id)  
  Rows Removed by Join Filter: 1637515  
  -> Seq Scan on trip (cost=0.00..17997.49 rows=58107 width=80) (actual  
      time=0.039..179.165 rows=58161 loops=1)  
      Filter: (bike_id < 200)  
      Rows Removed by Filter: 611798  
  -> Materialize (cost=0.00..2.05 rows=70 width=14) (actual time  
      =0.000..0.002 rows=29 loops=58161)  
      -> Seq Scan on station (cost=0.00..1.70 rows=70 width=14) (  
          actual time=0.016..0.048 rows=70 loops=1)  
Planning Time: 0.460 ms  
Execution Time: 599.250 ms
```

Based on the execution plan:

- | | |
|-------------------------------------------------------|--------------------------|
| i. Which join method was used? | <u>Nested loops join</u> |
| ii. What was the estimated cost (in arbitrary units)? | <u>78164.76</u> |
| iii. What was the total runtime (in ms)? | <u>599.250</u> |

(e) Execute these commands to re-enable the different joins.

```
RESET enable_mergejoin;  
RESET enable_hashjoin;  
RESET enable_indexscan;  
RESET enable_bitmapscan;
```